

# Introduction to R

Anders Stockmarr

Course developers: Anders Stockmarr, Elisabeth Wreford Andersen

DTU Department of Applied Mathematics and Computer Science  
Section for Statistics and Data Analysis  
Technical University of Denmark  
`anst@dtu.dk`

January 3rd, 2025

# Plan for today

- Introduction to R
- Data management
- Loops
- Graphics

# Outline

- 1 Introduction to R
- 2 Importing Data to R
- 3 Description of Data
- 4 Modifying Data
- 5 Loops and Flow Controls in R
- 6 Saving Your Work
- 7 Graphics
  - Histogram
  - Box plot
  - Scatter Plot
  - Line plot

# Introduction to R

- **R** is a programming *language* and a programming *environment*.
- It is Free! Developed by users under a GNU license.
- Runs on a variety of platforms including Windows, Unix and MacOS. You can even get it for Android.
- Allows for fast implementation of new methods by user demand through packages.
- **R** has state-of-the-art graphics capabilities.

# Advantages of R

Frank Harrel (my highlighting):

- "One point that hasn't been made very explicitly is one of the greatest advantages of **R**:

**Getting your work done better and in less time.**

Hundreds of companies hire a multitude of SAS programmers to write code in an archaic language, the SAS macro language. I believe there is a real cost savings from R because of its value as a **data analysis**, **data manipulation**, and **graphics** environment. Instead of programming using an indirect syntax manipulation environment (SAS macros), in R you can program in a dynamic data-sensitive framework".

That was more than 10 years ago. Things have progressed since...

# Base R

- Base **R** and most **R** packages are available for download at the Comprehensive R Archive Network (CRAN).
- <http://www.cran.r-project.org>
- Base **R** includes basic data management, analysis and graphics tools.
- For non-specialized tasks, Base **R** is all you need.
- Specialized tasks may be handled by *packages*.
- We will download, install and use many packages during the course.
- Packages are not all very well-documented (depends on the contributor).
- Want to be sure about what your program does?
  - Use well-established packages only;
  - or write your own code.

# RStudio

- You can work directly in **R**.
- Many prefer another front end (GUI, **G**raphical **U**ser **I**nterface).
- We will use **RStudio**.
- Download from <http://www.posit.co/>

# RStudio

- The GUI **RStudio** has 4 windows.
- One for writing the commands (the "script").
  - Use script for reproducibility.
- One for results and interactive use.
- One for plots, help and packages.
- One showing which objects are resident in the **R** memory.



## RStudio

The screenshot displays the RStudio IDE interface. The main editor window shows a script with R code for calculating the probability of a specific outcome in a binomial distribution. The console is empty. The Environment pane on the right lists several data objects, including spatial data frames and matrices. The sidebar on the right contains links to R resources, manuals, and a search engine.

**Script Editor:**

```

354 NS <- 100
355 NS <- 1:10
356 NS <- x
357 NS <- 0
358
359 % her skal svar skrives på følgende form. Rigtigt svar angives med *.
360 \begin{ttemize}
361 \item [1*] \Bbox { 13:  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$  }
362 \item [2] \Bbox { 9 :  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$  }
363 \item [3] \Bbox { 10 :  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$  }
364 \item [4] \Bbox { 7 :  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$  }
365 \item [5] \Bbox { 18 :  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$  }
366 \end{ttemize}
367
368 % her kan facit skrives
369 \begin{fact}
370 1. Sandsynligheden for ikke at vinde med 5 terninger er  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$ , og denne sandsynlighed skal være under 10%.
371 Altså  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$ , og dermed
372  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$ 
373 \end{fact}
374
375 % her kan facit skrives
376 \begin{fact}
377 1. Sandsynligheden for ikke at vinde med 5 terninger er  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$ , og denne sandsynlighed skal være under 10%.
378 Altså  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$ , og dermed
379  $\frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{10}$ 
380 \end{fact}

```

**Environment Pane:**

Object	Class	Size
coasts.spdf	Formal class "SpatialLinesDataFrame"	
countries.spdf	Formal class "SpatialPolygonsDataFrame"	
Europe	Large SpatialPolygonsDataFrame (68 elements, 931.2 kb)	
frame.spdf	Formal class "SpatialPolygonsDataFrame"	
graticule.spdf	Formal class "SpatialLines"	
my.center	1 obs. of 2 variables	
nuts0.df	34 obs. of 10 variables	
nuts0.spdf	Formal class "SpatialPolygonsDataFrame"	
nuts1.df	111 obs. of 10 variables	
nuts1.spdf	Large SpatialPolygonsDataFrame (111 elements, 536.0 kb)	
nuts2.df	310 obs. of 10 variables	
nuts2.spdf	Large SpatialPolygonsDataFrame (310 elements, 1.2 Mb)	
nuts3.df	1448 obs. of 7 variables	
nuts3.spdf	Large SpatialPolygonsDataFrame (1448 elements, 4.8 Mb)	

**Resources:**

- R Resources:** Learning R Online, CRAN Task Views, R on StackOverflow, Getting Help with R.
- RStudio:** RStudio IDE Support, RStudio Community Forum, RStudio Cheat Sheets, RStudio Tip of the Day, RStudio Packages, RStudio Products.
- Manuals:** An Introduction to R, Writing R Extensions, R Data Import/Export.
- Reference:** The R Language Definition, R Installation and Administration, R Internals.
- Packages:** Search Engine & Keywords.
- Miscellaneous Material:** About R, License, NEWS, Authors, FAQ, User Manuals, Resources, Thanks, Technical papers.

# R as a calculator

```
2+2
```

```
[1] 4
```

```
(2*5)+(12/3)-(2^3)
```

```
[1] 6
```

```
exp(log(1))
```

```
[1] 1
```

```
sqrt(25)
```

```
[1] 5
```

```
log(2*2)
```

```
[1] 1.3863
```

```
log(2)+log(2)
```

```
[1] 1.3863
```

# Writing commands in R

- Commands are separated by either a new line or ;
- **R** is case sensitive: id is a different name than ID.
- The character `#` at the beginning of a line shows that the text in this line is a comment. I.e. the text is not executed.
- Help can be found on the internet or in **R** by writing `?` followed by the function you want to help about:

```
?plot
```

- or, in [RStudio](#), highlight the expression and press F1.

# Objects in R

- Both data and output from analyses are stored as **objects** (if stored);
- Some times, output is just displayed on the screen, and you need to *assign* the object to an identifier to keep it (see below).
- In fact, everything in the **R** memory is stored in **objects**.
- An object could be a vector, a matrix or a data frame.
- Values are assigned to objects using the assignment operator `<-`
- The operator `=` also works, but it is **not recommended** as it may confuse assignments with default values and logical expressions.
- We can see the objects of the current **R** session memory in **RStudio**, or by using the function `ls()`

```
a <- 2+5
```

```
A <- 10
```

```
ls()
```

```
[1] "a" "A"
```

## Generating a sequence

- Specify the first and last values separated by a colon.
- Otherwise use `seq()`

```
0:10
```

```
[1] 0 1 2 3 4 5 6 7 8 9 10
```

```
15:5
```

```
[1] 15 14 13 12 11 10 9 8 7 6 5
```

```
seq(from = 0, to = 1.2, by = 0.1)
```

```
[1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0 1.1 1.2
```

```
x <- seq(from = 0, to = 1.5, length = 11); x
```

```
[1] 0.00 0.15 0.30 0.45 0.60 0.75 0.90 1.05 1.20 1.35 1.50
```

# Generating repeats using rep()

```
rep(8, 5)
```

```
[1] 8 8 8 8 8
```

```
rep(1:4, each = 2)
```

```
[1] 1 1 2 2 3 3 4 4
```

```
rep(1:4, each = 2, times = 3)
```

```
[1] 1 1 2 2 3 3 4 4 1 1 2 2 3 3 4 4 1 1 2 2 3 3 4 4
```

# Functions in R

We assign a simple function to the identifier f:

```
>f<-function(x){x^2}; f(2)  
[1] 4
```

A function of two variables:

```
>f<-function(x,pow){x^pow}; f(2,2)  
[1] 4
```

A function with a default value:

```
>f<-function(x,pow=2){x^pow}; f(2,2); f(2);f(2,3)  
[1] 4  
[1] 4  
[1] 8
```

# Functions in R

We have already used many functions with and without default values:

- `"+"(2,2)`
- `sqrt(25)`
- `log(2)`
- `ls()`
- `":"(0,10)`
- `seq(from=0.1,to=1.2,by=0.1)`
- `rep(1:4,each=2,time=3)`

Many applications in **R** are built up as functions. You can see default arguments in the help files. Example: `log`.



# Data structures in R: Singles

- Logical, e.g:

```
> TRUE
```

```
[1] TRUE
```

```
> 1==2
```

```
[1] FALSE
```

- Single numbers, e.g:

```
> 1
```

```
[1] 1
```

```
> 1.2
```

```
[1] 1.2
```

- Character, e.g:

```
> "5"
```

```
[1] "5"
```

```
> "abc"
```

```
[1] "abc"
```

## Data structures in R: Vectors

Constructed via the concatenate function `c()`.

- Vector of numbers, e.g:

```
> c(1,1.2,pi,exp(1))  
[1] 1.000000 1.200000 3.141593 2.718282
```

- We can have vectors of other things too, e.g:

```
> c(TRUE,1==2)  
[1] TRUE FALSE  
> c("a","ab","abc")  
[1] "a" "ab" "abc"
```

- But not combinations, e.g:

```
> c("a",5,1==2)  
[1] "a" "5" "FALSE"
```

Note that R just turned everything into characters!

## Data structures in R: Matrices

- Columns of same type and same length:

```
> matrix(c(1,2,3,4,5,6)+pi,nrow=2)
```

```
[,1] [,2] [,3]
```

```
[1,] 4.141593 6.141593 8.141593
```

```
[2,] 5.141593 7.141593 9.141593
```

```
> matrix(c(1,2,3,4,5,6)+pi,nrow=2)<6
```

```
[,1] [,2] [,3]
```

```
[1,] TRUE FALSE FALSE
```

```
[2,] TRUE FALSE FALSE
```

# Data structures in R: Data frames

- Same length of columns but different types; spread-sheet data.
- Created from reading in data from external files;
- or by using the function `data.frame()` on a set of vectors.

```
> data.frame(treatment=c("active","active","placebo"),
+ bp=c(80,85,90))
```

```
  treatment  bp
1   active  80
2   active  85
3 placebo  90
```

- Compare to a matrix created with the `cbind()` command:

```
> cbind(treatment=c("active","active","placebo"),bp=c(80,85,90))
      treatment bp
[1,] "active"  80
[2,] "active"  85
[3,] "placebo" 90
```

# Data structures in R: Lists

- Different length of columns and different types.
- Most general object type.

```
> list(a=1,b="abc",c=c(1,2,3),d=list(e=matrix(1:4,2), f=function(x){x^2}))
$a
[1] 1
$b
[1] "abc"
$c
[1] 1 2 3
$d
$d$e
      [,1] [,2]
[1,]    1    3
[2,]    2    4
$d$f
function (x)
{
    x^2
}
```

- The objects returned from many of the built-in functions in R are fairly complicated lists.

# Importing Data to R

- The easiest is to use data saved as text files.
- Usually values in text files are separated, or delimited, by tabs or commas.
- First tell **R** where you want to find your data using the command `setwd()`.
- Check that all went according to plan with `getwd()`.

```
setwd("C:/Users/ANST/Teaching/02935 January 2025")  
getwd()  
  
[1] "C:/Users/ANST/Teaching/02935 January 2025"
```

# Importing Data to R

- The function `read.table()` can be used to read data saved as text.
- Wrappers: `read.csv()`, `read.csv2()` and `read.delim()`.
- Notice the option `sep = .`
- We are assigning the loaded data to objects.
- If you have an Excel sheet, then save as text.

```
Births.tab <- read.table("Data/Births.txt", header = TRUE, sep = "\t")
```

```
# SHORT FORM TO READ TAB SEPARATED
```

```
Births.delim <- read.delim("Data/Births.txt")
```

```
# SHORT FORM TO READ ; SEPARATED
```

```
Births.csv2 <- read.csv2("Data/Births.csv")
```

```
# SHORT FORM TO READ , SEPARATED
```

```
Births.csv1 <- read.csv("Data/Births.csv1")
```

- 
- The screenshot displays the RStudio interface with the following components:
- File Editor (Top):** Contains the R script with the following code:
 

```
## Create a data frame with 1000 rows and 10 columns
set.seed(12345)
df <- data.frame(
  gender = sample(c("M", "F"), 1000, replace = TRUE),
  height = runif(1000, min = 150, max = 200),
  weight = runif(1000, min = 50, max = 150),
  age = runif(1000, min = 18, max = 65),
  bmi = runif(1000, min = 18, max = 30),
  blood_pressure = runif(1000, min = 90, max = 180),
  heart_rate = runif(1000, min = 60, max = 100),
  cholesterol = runif(1000, min = 100, max = 300),
  glucose = runif(1000, min = 70, max = 150),
  insulin = runif(1000, min = 0.1, max = 1.0)
)
```
  - Environment (Top Right):** Shows the data frame 'df' with 1000 observations and 10 variables. The variables are: gender (factor), height (numeric), weight (numeric), age (numeric), bmi (numeric), blood\_pressure (numeric), heart\_rate (numeric), cholesterol (numeric), glucose (numeric), and insulin (numeric).
  - Console (Bottom):** Shows the execution of the script, including the creation of the data frame and the calculation of the BMI for each observation.
  - Source (Bottom Left):** Shows the R script code, including the creation of the data frame and the calculation of the BMI for each observation.



# Importing Data using RStudio

- In the Objects Window, click "Import Dataset"

The screenshot shows the RStudio interface with the 'Import Dataset' dialog box open. The 'Input File' tab is selected, showing a file named 'births.csv' selected. The 'Data Frame' tab shows a preview of the data with columns: id, height, weight, gestwks, gestcns, mategs, and sex. The R console shows the following code:

```
> getwd()
[1] "C:/ANST/Undervisning/Karser/ghd/kursus 1 basal statistik 2jan2019/Introdag"
> births.tab <- read.table("births.txt", header = TRUE, sep = "\t")
> births.delim <- read.delim("births.txt")
> births.csv2 <- read.csv2("births.csv")
> births.csv <- read.csv("births.csv")
> cdc <- read.csv("C:/ANST/Undervisning/Karser/ghd/kursus 1 basal statistik 2jan2019/Introdag/cdc.csv")
> view(cdc)
```

# Importing Data from other stats programs

- We can read data from a series of other statistical software packages using the package `foreign`. A similar package is `Hmisc`, but we will stick to `foreign`.

```
# INSTALL AN EXTRA PACKAGE  
install.packages("foreign")
```

```
# ACTIVATE THE PACKAGE  
library("foreign")
```

```
SPSS_Data <- read.spss("Data/SPSS_Data.sav", to.data.frame = TRUE)
```

```
Stata_Data <- read.dta("Data/string.dta")
```

- For SAS data not in XPORT format, use the `sas7bdat` package.

# Look at Your Data

There are several ways to look at the data (or parts of the data).

```
# FIRST FEW OBSERVATIONS
```

```
head(Births.tab)
```

	id	bweight	lowbw	gestwks	preterm	matage	hyp	sex	sexalph
1	1	2974	0	38.52	0	34	0	2	female
2	2	3270	0	NA	NA	30	0	1	male
3	3	2620	0	38.15	0	35	0	2	female
4	4	3751	0	39.80	0	31	0	1	male
5	5	3200	0	38.89	0	33	1	1	male
6	6	3673	0	40.97	0	33	0	2	female

# Look at Your Data

```
# LAST FEW OBSERVATIONS
```

```
tail(Births.tab)
```

	id	bweight	lowbw	gestwks	preterm	matage	hyp	sex	sexalph
495	495	2968	0	41.01	0	34	0	1	male
496	496	2852	0	38.45	0	28	0	2	female
497	497	3187	0	38.03	0	38	1	1	male
498	498	3054	0	38.50	0	26	0	2	female
499	499	3178	0	39.92	0	31	0	2	female
500	500	2918	0	37.97	0	31	0	1	male

```
# VARIABLE NAMES
```

```
names(Births.tab)
```

```
[1] "id"          "bweight" "lowbw"    "gestwks" "preterm" "matage"  "hyp"
[8] "sex"         "sexalph"
```

```
# VIEW THE DATA IN A NEW WINDOW; OFTEN THE HEAD WILL DO THOUGH
```

```
View(Births.tab)
```

# Missing values

- In real life examples it is very common to have missing values.
- In **R** missing values are coded as **NA** (not available).
- In your Excel file leave missing values blank, do not set them to 99 or 999.

	id	bweight	lowbw	gestwks	preterm	matage	hyp	sex	sexalph
1	1	2974	0	38.52	0	34	0	2	female
2	2	3270	0	NA	NA	30	0	1	male

# Accessing Observations

- Data are (usually) stored in a data frame object.
- Observations are the rows.
- Variables, either numerical or categorical, are the columns.
- We can access individual rows, columns and cells in the data frame.
- For this, we use the bracket operator: `object[row, column]`.

# Accessing Observations

```
# A SINGLE CELL
```

```
Births.tab[345, 4]
```

```
[1] 38.55
```

```
# LEAVING OUT A COLUMN NUMBER INDICATES THAT ALL COLUMNS  
# ARE CHOSEN. HERE ALL COLUMNS IN ROW 224
```

```
Births.tab[224 , ]
```

	id	bweight	lowbw	gestwks	preterm	matage	hyp	sex	sexalph
224	224	3216	0	39.94	0	38	1	1	male

## Accessing Observations

```
# LEAVING OUT A ROW NUMBER INDICATES THAT ALL ROWS ARE CHOSEN
# HERE ALL ROWS IN COLUMN 5
Births.tab[,5]
```

[illegible]



# Accessing Observations

```
# USE RANGES, ROWS 15 TO 18 COLUMNS 1 TO 4  
Births.tab[15:18, 1:4]
```

	id	bweight	lowbw	gestwks
15	15	3662	0	39.23
16	16	3035	0	38.96
17	17	3351	0	39.35
18	18	3804	0	38.99

## Accessing Observations

Variables can be accessed directly using their name, either with the `$` operator (`object$variable`), the name (`object[, "variable"]`), or the column number (`object[, k]`).

```
# GET THE BIRTH WEIGHT FOR CHILD 26 TO 36
```

```
Births.tab$bweight[26:36]
```

```
[1] 3585 3798 3164 3739 1780 4022 3942 2887 2391 3911 3509
```

```
Births.tab[26:36, "bweight"]
```

```
[1] 3585 3798 3164 3739 1780 4022 3942 2887 2391 3911 3509
```

```
Births.tab[26:36, 2]
```

```
[1] 3585 3798 3164 3739 1780 4022 3942 2887 2391 3911 3509
```

## Subsetting using the `c()` function

- The concatenate function `c()` concatenates the arguments into a vector. It can be used for many things; one is to access non-sequential rows and columns from a data frame.

```
# GET COLUMNS 2, 5, 7, 8, 9 FOR ROW 33
```

```
Births.tab[33, c(2, 5, 7:9)]
```

```
      bweight preterm hyp sex sexalph
33      2887         0  0  1      male
```

```
# GET bweight, preterm and sexalph FOR ROW 71
```

```
Births.tab[71, c("bweight", "preterm", "sexalph")]
```

```
      bweight preterm sexalph
71      3189         0      male
```

# Variable Names

If we want to change the variable names we can use `names()`.

```
# NEW VARIABLE NAMES
names(Births.tab) <- c("ID", "Bweight", "LowBW", "GestWks",
                      "Preterm", "Matage", "Hyp", "Sex", "Sexalph")

# JUST THE FIRST VARIABLE NAME
names(Births.tab)[1] <- c("ID_new")

# CHECK HOW IT WENT
names(Births.tab)

[1] "ID_new"  "Bweight" "LowBW"   "GestWks" "Preterm" "Matage"  "Hyp"
[8] "Sex"     "Sexalph"

# RESETTING
names(Births.tab)[1] <- c("ID")
```

## Saving/Exporting data

- We can save the data to a textfile using either `write.table()` for a tab separated file, or `write.csv()/write.csv2()` for a comma/semicolon separated file (with "." and "," as decimal point, respectively).

```
write.table(Births.tab, file = "Birth_new.txt",  
            sep = "\t", na = ".", row.names= FALSE)  
  
write.csv2(Births.tab, file = "Birth_new.csv")
```

## Exercise: Protein Consumption

- Open [RStudio](#) and set the working directory to where you want to keep the data for the course today.
- Import the data Protein.xlsx to **R**.
- Look at the data.
- What is the protein consumption in Denmark?
- Look at the protein consumption from red meat alone.
- Look at the protein consumption in Denmark, Norway, Sweden from red meat, white meat and eggs.
- Rename the variables “RedMeat” and “WhiteMeat” to “Red” and “White”.
- Save the new version of the protein data as a tab delimited text file “Protein2.txt”.

# Description of Data

We are still looking at the data set with birth weights for 500 children. Using the function `str()` we can see a description of what our data frame contains (the structure).

```
str(Births.tab)
```

```
'data.frame': 500 obs. of 9 variables:
 $ id      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ bweight: int  2974 3270 2620 3751 3200 3673 3628 3773 3960 3405 ...
 $ lowbw   : int   0 0 0 0 0 0 0 0 0 0 ...
 $ gestwks: num   38.5 NA 38.2 39.8 38.9 ...
 $ preterm: int   0 NA 0 0 0 0 0 0 0 0 ...
 $ matage  : int   34 30 35 31 33 33 29 37 36 39 ...
 $ hyp     : int   0 0 0 0 1 0 0 0 0 0 ...
 $ sex     : int   2 1 2 1 1 2 2 1 2 1 ...
 $ sexalph: chr  "female" "male" "female" "male" ...
```

# Description of Data: Birth weights

- `str()` supplies all information about an object. **Much is redundant**, do not use `str()` for report writing but for your own overview only.
- We see that we have a data frame with 500 observations and 9 variables.
- Some are integers but “gestwks” is numeric.
- The variable “sexalph” is a character variable.
- Note that “sexalph” and “sex” describes the same thing. But R does not interpret the character values as group labels.
- Factor: Grouping with informative labels. We can convert “sexalph” to a factor using `as.factor()`.



# Description of Data: Birth weights

```
# TELL R THAT sexalph IS A FACTOR
```

```
Births.tab$sexalph <- as.factor(Births.tab$sexalph)
```

```
levels(Births.tab$sexalph)
```

```
[1] "female" "male"
```

```
str(Births.tab)
```

```
'data.frame': 500 obs. of 9 variables:
```

```
$ id      : int  1 2 3 4 5 6 7 8 9 10 ...
```

```
$ bweight: int  2974 3270 2620 3751 3200 3673 3628 3773 3960 3405 ...
```

```
$ lowbw   : int  0 0 0 0 0 0 0 0 0 0 ...
```

```
$ gestwks: num  38.5 NA 38.2 39.8 38.9 ...
```

```
$ preterm: int  0 NA 0 0 0 0 0 0 0 0 ...
```

```
$ matage  : int  34 30 35 31 33 33 29 37 36 39 ...
```

```
$ hyp     : int  0 0 0 0 1 0 0 0 0 0 ...
```

```
$ sex     : int  2 1 2 1 1 2 2 1 2 1 ...
```

```
$ sexalph: Factor w/ 2 levels "female","male": 1 2 1 2 2 1 1 2 1 2 ...
```

# Description of Data: Birth weights

```
# TELL R THAT sex IS A FACTOR WITH SPECIFIC LEVELS
```

```
Births.tab$sex <- factor(Births.tab$sex, labels = c("Male", "Female"))
levels(Births.tab$sex)
```

```
[1] "Male" "Female"
```

```
str(Births.tab)
```

```
'data.frame': 500 obs. of 9 variables:
```

```
$ id      : int  1 2 3 4 5 6 7 8 9 10 ...
```

```
$ bweight: int  2974 3270 2620 3751 3200 3673 3628 3773 3960 3405 ...
```

```
$ lowbw   : int  0 0 0 0 0 0 0 0 0 0 ...
```

```
$ gestwks: num  38.5 NA 38.2 39.8 38.9 ...
```

```
$ preterm: int  0 NA 0 0 0 0 0 0 0 0 ...
```

```
$ matage  : int  34 30 35 31 33 33 29 37 36 39 ...
```

```
$ hyp     : int  0 0 0 0 1 0 0 0 0 0 ...
```

```
$ sex     : Factor w/ 2 levels "Male","Female": 2 1 2 1 1 2 2 1 2 1 ...
```

```
$ sexalph: Factor w/ 2 levels "female","male": 1 2 1 2 2 1 1 2 1 2 ...
```

# Descriptive Statistics

- There are some simple functions for summary statistics in **R**.
- Very common extractor functions are `mean()`, `sd()`, `median()`, `max()` and `min()`.

```
mean(Births.tab$bweight)
```

```
[1] 3136.9
```

```
sd(Births.tab$bweight)
```

```
[1] 637.45
```

```
median(Births.tab$bweight)
```

```
[1] 3188.5
```

```
max(Births.tab$bweight)
```

```
[1] 4553
```

```
min(Births.tab[, 2])
```

```
[1] 628
```

# The Summary Function

- The function `summary()` can be used with many objects in **R**.
- When used on a data frame we get all the main summary statistics.

```
# SUMMARY OF THE DATA FRAME
```

```
summary(Births.tab)
```

id		bweight	lowbw	gestwks	
Min.	: 1	Min. : 628	Min. :0.00	Min.	:24.7
1st Qu.:	126	1st Qu.:2862	1st Qu.:0.00	1st Qu.:	37.9
Median :	250	Median :3188	Median :0.00	Median :	39.1
Mean :	250	Mean :3137	Mean :0.12	Mean :	38.7
3rd Qu.:	375	3rd Qu.:3551	3rd Qu.:0.00	3rd Qu.:	40.1
Max.	:500	Max. :4553	Max. :1.00	Max.	:43.2
				NA's	:10

preterm		matage	hyp	sex	sexalph
Min.	:0.000	Min. :23	Min. :0.000	1:264	female:236
1st Qu.:	0.000	1st Qu.:31	1st Qu.:0.000	2:236	male :264
Median :	0.000	Median :34	Median :0.000		
Mean :	0.129	Mean :34	Mean :0.144		
3rd Qu.:	0.000	3rd Qu.:37	3rd Qu.:0.000		
Max.	:1.000	Max. :43	Max. :1.000		
NA's	:10				

# Summaries

- We may only want summaries for some of the data, e.g. babies with birth weight  $< 2900\text{g}$ .
- We subset the data and then summarize as before:

```
summary(Births.tab[Births.tab$bweight<2900,])
```

id	bweight	lowbw	gestwks
Min. : 3	Min. : 628	Min. : 0.000	Min. : 24.7
1st Qu.: 146	1st Qu.: 2120	1st Qu.: 0.000	1st Qu.: 35.5
Median : 254	Median : 2580	Median : 0.000	Median : 37.4
Mean : 250	Mean : 2355	Mean : 0.441	Mean : 36.6
3rd Qu.: 359	3rd Qu.: 2741	3rd Qu.: 1.000	3rd Qu.: 38.5
Max. : 496	Max. : 2894	Max. : 1.000	Max. : 41.4
			NA's : 2
preterm	matage	hyp	sex      sexalph
Min. : 0.000	Min. : 24	Min. : 0.000	1:63   female:73
1st Qu.: 0.000	1st Qu.: 31	1st Qu.: 0.000	2:73   male :63
Median : 0.000	Median : 34	Median : 0.000	
Mean : 0.403	Mean : 34	Mean : 0.243	
3rd Qu.: 1.000	3rd Qu.: 37	3rd Qu.: 0.000	
Max. : 1.000	Max. : 43	Max. : 1.000	
NA's : 2			

## Group Summaries

- Data may be separated by groups.
- Suppose that we want to calculate the mean birth weight for boys and girls (many ways to do this).
- We will use the `tapply()` function to apply the `mean` function to the two levels of “sexalph”.
- `tapply(<variable to summarize>, <variable to group by>, <function to use>)`.

```
# MEAN BIRTH WEIGHT FOR BOYS AND GIRLS
```


```
tapply(Births.tab$bweight, Births.tab$sexalph, mean)
```

```
female      male  
3032.831 3229.902
```

# Histogram

Often it is easier to get an impression of a distribution using plots. Histograms are typically used for continuous variables.


```
hist(Births.tab$bweight, main = "Title", xlab = "Birth weight (g)" )
```



# Histogram

Often it is easier to get an impression of a distribution using plots. Histograms are typically used for continuous variables. Here with a box on.

```
hist(Births.tab$bweight, main = "Title",  
      xlab = "Birth weight (g)")  
box()
```






# Boxplot

Boxplots show the median, upper, lower quartiles and potentially extreme values.

```
boxplot(Births.tab$bweight, xlab = "Birth weight (g)")
```



## Exercise: Descriptive Statistics 1

- Import the dataset `cdc.csv` into **R**.
- Look at the data.
- Describe the structure of the data.
- Change the variable “`smoke100`” to a factor.
- Summarize the variables in the dataset.
- Calculate the mean age by general health.
- Draw a histogram of age.
- Draw a boxplot of age.

# Modifying Data

We will concentrate on how to modify and rearrange our data.

- Data can be sorted with the `order` function.
- `order` can sort the `Birth.tab` data by “sex“, and then by “bweight“.
- The `order` function returns a vector of sorted indices, which we apply to the rows of the unsorted data frame to get a sorted version.

```
Birth_sort <- Births.tab[order(Births.tab$sex, Births.tab$bweight), ]
```

```
head(Birth_sort)
```

	id	bweight	lowbw	gestwks	preterm	matage	hyp	sex	sexalph
253	253	693	1	30.71	1	34	1	1	male
226	226	981	1	27.99	1	29	1	1	male
181	181	1019	1	28.04	1	31	1	1	male
22	22	1203	1	32.80	1	39	0	1	male
312	312	1500	1	35.27	1	34	0	1	male
313	313	1595	1	30.52	1	33	1	1	male

## Creating new variables and deleting old

New variables can be added to a data frame.

```
# ADD A VARIABLE TO DATA FRAME  
Births.tab$log_bweight <- log(Births.tab$bweight)
```

Columns can be dropped from a data frame (log birth weight is column 10):

```
Births.tab <- Births.tab[ , -10]
```

```
# CREATE A VARIABLE AS A SEPARATE OBJECT  
log_bweight <- log(Births.tab$bweight)
```

Objects can be removed from the R memory (cleaning up):

```
rm(log_bweight)
```

## Grouping the values of a variable using cut

You might want to group a continuous variable e.g. mother's age (matage) into the groups: ]20-30], ]30-35], ]35-40], ]40-45]:

```
Births.tab$agegrp <- cut(Births.tab$matage,
                        breaks = c(20, 30, 35, 40, 45))
summary(Births.tab[, c("matage", "agegrp")])
```

	matage	agegrp
Min.	:23	(20,30]: 99
1st Qu.:	:31	(30,35]: 215
Median	:34	(35,40]: 174
Mean	:34	(40,45]: 12
3rd Qu.:	:37	
Max.	:43	

## Creating new variables: RowSums

- Often we want to form new variables from other variables.
- For example, we might want to calculate a total score from sub scores.
- We can sum variables using `rowSums`. Related functions are: `rowMeans`, `colSums`, `colMeans`.
- Notice the option `na.rm`.
- If we take a row sum where one of the values is missing then the row sum is set to missing `na.rm= FALSE`.
- If we want to ignore missing values and calculate a sum of the non missing then `na.rm= TRUE`.
- `rowSums`, `rowMeans`, `colSums` and `colMeans` are wrappers of `sapply`, ie. t.ex. `colMeans(x)` is the same as `sapply(x,mean)`. `sapply` can be used with many other functions.

# Creating new variables: RowSums

```
# NEW VARIABLE SCORE SUMMING PRETERM, LOWBW AND HYP
Births.tab$score <- rowSums(Births.tab[,c(3,5,7)], na.rm = FALSE)

#REMOVE MISSING
Births.tab$scoreRM <- rowSums(Births.tab[,c(3,5,7)], na.rm = TRUE)

head(Births.tab)
```

	id	bweight	lowbw	gestwks	preterm	matage	hyp	sex	sexalph	score	scoreRM
1	1	2974	0	38.52	0	34	0	2	female	0	0
2	2	3270	0	NA	NA	30	0	1	male	NA	0
3	3	2620	0	38.15	0	35	0	2	female	0	0
4	4	3751	0	39.80	0	31	0	1	male	0	0
5	5	3200	0	38.89	0	33	1	1	male	1	1
6	6	3673	0	40.97	0	33	0	2	female	0	0

## Exercise

- Import the data cdc.
- Height is in inches, weight and wt desire are in pounds. Generate new variables in cm and kg (use Google to find conversion factors).
- Make a factor with 4 *approximately* equally sized groups from the weight variable.
- Sort the data by gender and age.
- Calculate the average of the weight and desired weight for each subject.
- Calculate the mean weight and wt desire (in kg) for each level of "genhlth" with the `by` function.



## Split Data: Subset

- Sometimes we may need to split our data.
- In the Births data we may need to split the data into boys and girls.
- We can use the `subset()` function and assign the new data sets to separate **R** objects.
- Notice `==` (logical expression). We are not assigning a value to “sex”, but asking whether “sex is equal to 1”.

```
Births.Male <- subset(Births.tab, sex == 1)
Births.Female <- subset(Births.tab, sex == 2)
```

- Alternatively the bracket operator:

```
Births.Male <- Births.tab[Births.tab$sex == 1,]
Births.Female <- Births.tab[Births.tab$sex == 2,]
```

# Subset

- Often data sets come with a lot of variables and we only want to use a few.
- Similarly to the bracket operator `[` the function `subset()` can be used to select the variables we want.
- Notice the `select` option. This is needed to say that we want a subset of columns (on the previous slide it was rows).
- Notice that we do not need quotes in `select`.

```
# SELECT 3 VARIABLES
```

```
Births.new <- subset(Births.tab, select = c(id, bweight, sex))
```

## Aggregating data

- Sometimes we want to make a new dataframe as a summary of the original dataframe on the basis of factor levels.
- Below we want to make a new dataframe with the mean birthweight for combinations of preterm and sex.

```
PreSex <- aggregate(Births.tab$bweight,
                    by = list(Preterm = Births.tab$preterm,
                              Sex = Births.tab$sexalph), mean)
```

PreSex

	Preterm	Sex	x
1	0	female	3191.6
2	1	female	2020.8
3	0	male	3361.2
4	1	male	2321.6

## Add rows: `rbind()`

- Data are collected for subgroups of subjects and saved in separate objects.
- The separate objects are appended (stacked) to create a single object.
- This will give an error message if the number of columns differs.

```
# APPEND
```

```
Births.Both <- rbind(Births.Male, Births.Female)
```

```
dim(Births.Both)
```

```
[1] 500  11
```

```
dim(Births.Male)
```

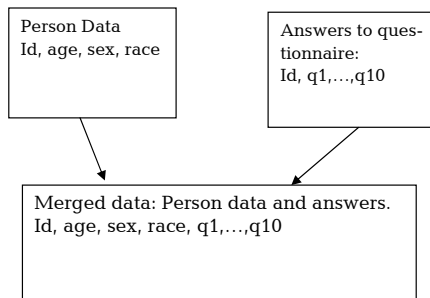
```
[1] 264  11
```

```
dim(Births.Female)
```

```
[1] 236  11
```

## Add variables: merge()

Often you have data in several data sets and want to combine the data sets by merging using one or more variables as *key variables*. Adding variables to a master data set.



# Merge

We have two data sets with a key variable "id". One with background information and one set with blood pressure measurements.

```
agesex <- read.delim("agesex.txt")
```

```
agesex
```

	id	age	sex
1	99	43	m
2	100	47	f
3	101	NA	f
4	102	67	m


```
bp <- read.delim("bp.txt")
```

```
bp
```

	id	visit	bp
1	100	1	180
2	100	2	160
3	100	3	155
4	101	1	160
5	102	1	120
6	102	2	140
7	103	1	135

## 4 Different Merges

- In the `merge` function we will look at 4 of the options.
- We have `merge(x, y, by = "key variable", all = TRUE, < all = FALSE, all.x = TRUE, all.y = FALSE > )`.
- Here `x` and `y` are data frames



# Merging all=FALSE

```
merge_small <- merge(agesex, bp, by = "id", all = FALSE)
merge_small
```

	id	age	sex	visit	bp
1	100	47	f	1	180
2	100	47	f	2	160
3	100	47	f	3	155
4	101	NA	f	1	160
5	102	67	m	1	120
6	102	67	m	2	140



# Merging all=TRUE

```
merge_large <- merge(agesex, bp, by = "id", all = TRUE)
merge_large
```

	id	age	sex	visit	bp
1	99	43	m	NA	NA
2	100	47	f	1	180
3	100	47	f	2	160
4	100	47	f	3	155
5	101	NA	f	1	160
6	102	67	m	1	120
7	102	67	m	2	140
8	103	NA	<NA>	1	135

## Merging all.x=TRUE

```
merge_x <- merge(agesex, bp, by = "id", all.x = TRUE)  
merge_x
```

	id	age	sex	visit	bp
1	99	43	m	NA	NA
2	100	47	f	1	180
3	100	47	f	2	160
4	100	47	f	3	155
5	101	NA	f	1	160
6	102	67	m	1	120
7	102	67	m	2	140

## Merging all.y=TRUE

```
merge_y <- merge(agesex, bp, by = "id", all.y = TRUE)  
merge_y
```

	id	age	sex	visit	bp
1	100	47	f	1	180
2	100	47	f	2	160
3	100	47	f	3	155
4	101	NA	f	1	160
5	102	67	m	1	120
6	102	67	m	2	140
7	103	NA	<NA>	1	135

# Counting the Missing Observations: The `is.na()` and `sum()` functions

- Suppose that we want to count the number of missing observations.
- The function `is.na` returns a logical vector that is **TRUE** when a value is missing and **FALSE** otherwise.

```
is.na(merge_y$sex)

[1] FALSE FALSE FALSE FALSE FALSE FALSE TRUE
```

```
#COUNT MISSING FOR ONE VARIABLE
```

```
sum(is.na(merge_y$sex))
```

```
[1] 1
```

```
#COUNT FOR DATA FRAME
```

```
colSums(is.na(merge_y))
```

```
id    age    sex visit    bp
0      2      1     0     0
```

# Loops in R

- In **R**, the **for loop** is used to perform a repetitive task for each element in a set.
- Example:
  - Given a set of integers 1:3:
  - Let a variable  $i$  run through the set and print  $i + i$ :

```
for(i in 1:3) {  
  cat(i, "+", i, "=", i+i, "\n")  
}
```

- Output:

1 + 1 = 2

2 + 2 = 4

3 + 3 = 6

# Flow Control: if and if else Statements

- if statement:

```
for(i in 1:3){
  if (i==2) cat("This index is even:","\n")
  cat(i,"\n")
}
```

1

This index is even:

2

3

- if else statement:

```
for(i in 1:3){
  if (i==2) cat("The index is 2","\n") else
    cat("The index is not 2","\n")
}
```


The index is not 2

The index is 2


The index is not 2

# Flow Charts for if and if else Statements

Enter if statement



Enter if else statement



# while and repeat Loops

The **while** loop:

- `while(condition) expression`

The **repeat** loop:

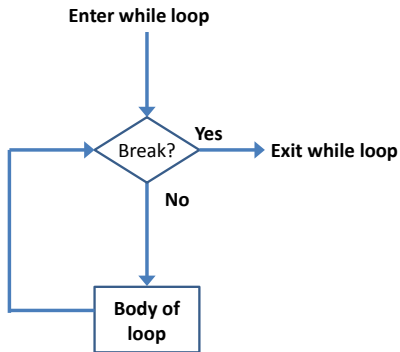

- `repeat expr`

The repeat loop has to be exited manually. Flow controls:

- **next:** Halts the the current iteration and advances to the next immediately;
- **break:** Exits the loop.



# Flow Charts for the while and repeat loops



## while Loop Example

### Storing of machine parts

```
k<-0 # number of big parts (>2)
y<-abs(rnorm(1000)) # simulated part size
i<-0 # index of parts
# loop:
while(k<3 & i<1000){
  i<-i+1
  temp<-y[i]
  k<-k+(temp>2)
}
i

[1] 42
```

## repeat Loop Example

### Selecting persons without blue or yellow eyes

```

eye.colors<-c("brown","blue","green","yellow","grey")
eyecolor<-data.frame(personId=1:100,color=
                      sample(eye.colors,100,rep=T))

i<-0
list.of.ids<-numeric(0) # patient ID list
#loop:
repeat {
  i<-i+1
  if(eyecolor$color[i]=="yellow" |
     eyecolor$color[i]=="blue") next
  list.of.ids<-c(list.of.ids,eyecolor$personId[i])
  if(i==100 | length(list.of.ids)==20) break
}
list.of.ids

```

```
[1] 5 6 7 9 10 11 12 14 15 18 19 20 21 22 23 24 25 28 29 30
```

# Saving your work

- Saving your script
- Saving your workspace

Always save your script - do it often if you work in [Rstudio](#).

- Reasons for saving your workspace:
  - Extensive data creations will be there next time you open your workspace.
  - Objects created 'on the fly' (not in your script) will be there.
- Reasons for not saving your workspace:
  - With a well-written script, you can recreate your analysis in seconds, unless you work with huge amounts of data.
  - Edited and saved data where editions have been forgotten may cause havoc on your results.
  - Left-over objects created for various purposes may enter your calculations unintentionally due to the structure of **R**'s search path.

# Saving your work

How to save your work:

- Script: Click on the script and press 'save' in [Rstudio](#) and the plain R GUI.
- Workspace: Click on the command prompt and press 'save'.  
Alternatively, use the [save.image\(\)](#) function
- Both: Accept when asked after terminating [Rstudio](#) or the plain R GUI.

## Exercise

- Data on exercise habits for the subjects in the cdc data set are found in the data `cdc_exer.csv`. Load and describe these data.
- Merge the cdc data set with the `cdc_exer` data set using `id` as the key variable. How many subjects in the study did not have exercise information?
- List the identifiers for the subjects without exercise information.
- Make a new dataframe with average age for combinations of gender and general health.

# Visualizing Data is Important


- Whenever we want to analyze data the first thing we should do is to **have a look** at it.
- How are the observations spread out?
- What are the most common values?
- Are there any unusual observations?
- Are there any relationships between variables?

This session will not tell you all about graphics in **R** but get you going.

# A Basic Histogram

- Common way to examine the distribution of a continuous variable.
- The range of the variable is by default divided into equal-width intervals (bins). Plots the number of observations in each bin (unless you specify otherwise).

```
hist(Births$bweight)
```






## Histogram with a few options

- Note that **R** automatically has created axis labels and a heading.
- To modify axis labels we set the options `xlab` and `ylab`.
- The heading is set in the option `main`.


```
hist(Births$bweight, xlab = "Birth weight (g)",  
     main = "Histogram of birth weight")
```



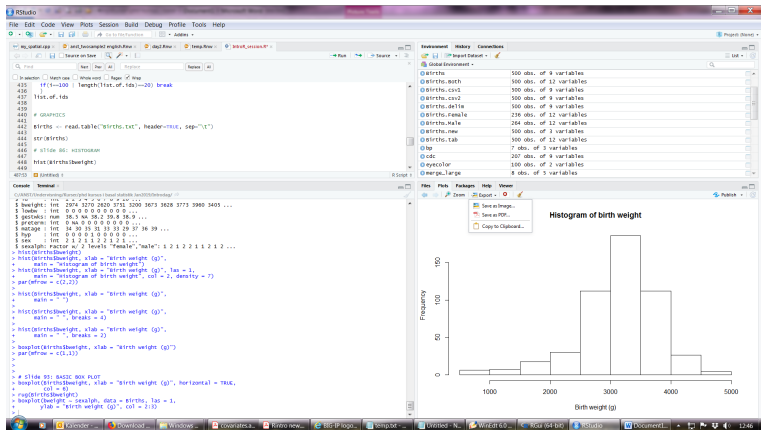
# Histogram with more options

- We could type `?hist` to find more options to customize the histogram.
- The available colours are coded as numbers or one can write `col = "red"`
- If we want shading we can try the `density` function.
- The angle of the numbers on the axes is set by the option `las`.

```
hist(Births$bweight, las = 1, main = "Histogram of birth weight",  
     col = 2, density = 7)
```



# How to get your plot from RStudio



# Writing to a graphics device

```
pdf("my.histogram.pdf")
hist(Births$bweight, las = 1, main = "Histogram of birth weight",
      col = 2, density = 7)
dev.off()
```

```
svg("my.histogram.svg")
hist(Births$bweight, las = 1, main = "Histogram of birth weight",
      col = 2, density = 7)
dev.off()
```

```
png("my.histogram.png")
hist(Births$bweight, las = 1, main = "Histogram of birth weight",
      col = 2, density = 7)
dev.off()
```

- Options can be specified; see the help files.





## Exercise

- Load the cdc data set.
- Make a histogram of the weight in kg.
- Add your own title and x-axis label.
- Try different colours and shadings.
- Copy your favorite histogram into a document in e.g. Word.

# A Basic Box Plot

- Box plots show some distributional properties very clearly.
- Box plots show a measure of the location (the median line).
- The spread of the distribution (the length of the box and whiskers).
- Skewness as asymmetry in the upper and lower parts of the box and whisker length.
- We use the function `boxplot(variable)`. Adding labels to the axes and colours is done as for `hist`.


# Histograms and a Box Plot



# A Basic Box Plot

- When describing data we can even add the observations to the plot.
- Notice the function `rug` shows the observations.

```
boxplot(Births$bweight, xlab = "Birth weight (g)", horizontal = TRUE,  
        col = 6)  
rug(Births$bweight)
```






## Box Plot for Groups

A useful feature is that we can make box plots for different groups next to each other for comparison. Notice the option `data = Births`.

```
# BOX PLOT FOR BOYS AND GIRLS
```

```
boxplot(bweight ~ sexalph, data = Births, las = 1,  
        ylab = "Birth weight (g)", col = 2:3)
```




## Box Plot for Groups

Set our own axis. Notice `xaxt = "n"`.

```
# BOX PLOT WHERE WE WANT TO MAKE OUR OWN AXIS
```

```
boxplot(bweight ~ sexalph, data = Births, las = 1,  
        ylab = "Birth weight (g)", col = c("red", "blue"), xaxt = "n")  
axis(1, at = c(1,2), labels = c('Girl', 'Boy'))
```



## Exercise


- Load the cdc data set.
- Make a box plot of the weight in kg showing individual observations. Add your own titles and colours.
- Make a box plot for each level of "genhlth", where each box is a different colour. How were the plots sorted?
- Make a new variable genhlth.num with 5 levels where 1 is "poor", 2 is "fair", 3 is "good", 4 is "very good" and 5 is "excellent".
- Make a new box plot for each level of genhlth.num with labels on the x-axis.
- Does this box plot show any pattern?

# The Basic Scatter Plot

- The scatter plot is the standard graph for examining the relationship between two continuous variables.
- The `plot(x,y)` function is used to create scatter plots. Where  $(x,y)$  are the points we want to plot.
- We will look at the relationship between car weight (lbs/1000) and miles per gallon for 32 cars.

```
plot(mtcars$wt, mtcars$mpg)
```


# The Basic Scatter Plot



# The Scatter Plot

- We can customize the scatter plot similar to before.
- The function `abline` adds a straight line to the plot.
- When we write `abline(lm(mpg ~ wt))` we get the best fitting line.


```
plot(mtcars$wt, mtcars$mpg, xlab = "Car weight (lbs/1000)",  
     ylab = "Miles per gallon", las = 1, pch = 19)  
abline(lm(mtcars$mpg ~ mtcars$wt), lty = 1, col = 3)
```



# abline

- The function `abline` can also add reference lines to a plot.
- A horizontal line, e.g. at 25 and 30 `abline(h = c(25, 30))`
- A vertical line, e.g. at 2 and 5 `abline(v = c(2, 5))`

```
plot(mtcars$wt, mtcars$mpg, xlab = "Car weight (lbs/1000)",  
     ylab = "Miles per gallon", las = 1, pch = 19)  
abline(h = c(25, 30), col = c("red", "magenta"), lty = 2)  
abline(v = c(2, 5), col = 4:5, lty = 3:4)
```




## Add a smoothed line

Perhaps we do not think the association is linear and try a nonparametric smoothed line.

```
plot(mtcars$wt, mtcars$mpg, xlab = "Car weight (lbs/1000)",  
     ylab = "Miles per gallon", las = 1)  
abline(lm(mtcars$mpg ~ mtcars$wt), lty = 2, col = 4)  
lines(lowess(mtcars$wt, mtcars$mpg), lty = 1, col = 2)
```



## Add a smoothed line




## Enhanced graph procedures: Scatter plot example from the "car" package

```
scatterplot(mpg ~ wt | cyl, data = mtcars, ylim = c(0,40),  
            xlab = "Car weight (lbs/1000)",  
            ylab = "Miles per gallon", las = 1,  
            legend = list(coords="topright"),  
            id = list(method="identify"),  
            boxplots = "xy", col=2:4)
```


- Here we want to plot miles per gallon versus weight for cars that have 4, 6 or 8 cylinders. We write this as `mpg ~ wt | cyl`.
- By default we get different colours for groups and both a linear and a smoothed line.
- A legend is included in the top right corner of the plot.
- The option `id = list(method="identify")` means that points can be identified by mouse clicks.
- Box plots of miles per gallon and weight included ("xy" option for both axes).
- More possibilities: `?scatterplot`.

# The resulting scatter plot




# A scatter plot matrix from the "car" package

```
scatterplotMatrix( ~ mpg + hp + drat + wt, data = mtcars,  
  smooth = list(col.smooth=2,col.spread=3,  
    lty.smooth=1),  
  regLine = FALSE,ellipse = TRUE)
```



# A scatter plot matrix with boxplots

```
scatterplotMatrix( ~ mpg + hp + drat + wt | cyl,
  diagonal = list(method="boxplot"),
  smooth = list(lty.smooth=1),
  regLine = FALSE, col = 2:6, data = mtcars )
```



## Exercise

- Load the cdc data set.
- Make a basic scatter plot of weight and desired weight with “nice” axes and labels.
- Add a smoothed lined to the plot.
- Make a scatter plot of weight and desired weight for each level of “genhlth” in separate plots, with straight and smoothed lines.
- Install the package “car”.
- Make a scatter plot of weight and desired weight for each level of “genhlth” in one plot.


## Exercise

- Load the Protein data set.
- Make a scatter plot matrix of “RedMeat”, “WhiteMeat”, “Eggs”, “Milk” and “Fish”. Can you see any patterns in the protein intake?
- Also make a scatter plot matrix with box plots in the diagonal.

# A Line Plot

Connecting points in a scatter plot from left to right. Here the growth of a tree. Notice the option `type = "b"` meaning points joined by lines.

```
par(mfrow=c(1,2)),  
plot(TreeA$age, TreeA$circumference, xlab = "Age (days)",  
      ylab = "Circumference (mm)", las = 1)  
plot(TreeA$age, TreeA$circumference, type = "b", xlab = "Age (days)",  
      ylab = "Circumference (mm)", las = 1)
```






## Difference between `plot()` and `lines()` functions

- We have seen both the `plot` and the `lines` functions.
- The `plot` function creates a new graph. It is a **high-level** plotting function.
- The `lines` function adds information to an existing graph but it cannot produce it's own graph. It is a **low-level** plotting function.
- A **high-level** plotting function can (often) be converted to a **low-level** plotting function with the option `ADD=TRUE`.
- Usually `lines` will be used after a **high-level** plotting function (such as `plot`) has produced a graph.


# A line plot and a legend

```
plot(TreeA$age, TreeA$circumference, type = "b", lty = 1,  
      xlab = "Age (days)", ylab = "Circumference (mm)", las = 1, col = 2)  
lines(TreeB$age, TreeB$circumference, type = "b", col = 3, lty = 2)  
legend(locator(1), # we will place it with a mouse click  
       legend = c("A", "B"), title = "Tree",  
       lty = 1:2, col = 2:3)
```



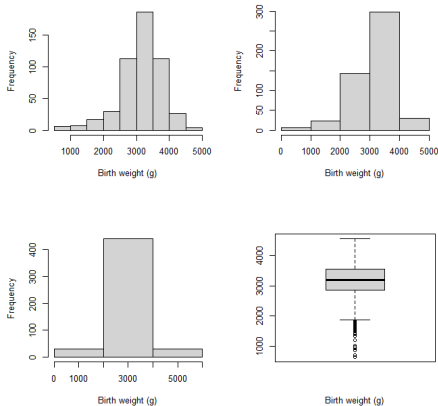
# Plot of growth of 5 trees

```
treedata<-Orange
plot(treedata$age,treedata$circumference,pch=" ",xlab = 'Age',
     ylab='Circumference',las=1)
for(i in 1:5){lines(treedata$age[treedata$Tree==i],
                    treedata$circumference[treedata$Tree==i],type="b",col=i+1)}
legend(x="topleft",paste("Tree #",1:5),pch=1,lty=1,col=2:6)
```



# Layout of several plots on one graph

Several plots on one graph:



Use the option `par(mfrow = c(2, 2))`, and back to one plot `par(mfrow = c(1, 1))`.

## Layout of several plots on one graph

The `layout()` function:

- indicate in matrix form which part of the plot area that you wish to belong to which graph

```
> layout.matrix<-matrix(c(1:3,rep(4,3)),nrow=2,ncol=3,byrow=T)
> layout.matrix
      [,1] [,2] [,3]
[1,]    1    2    3
[2,]    4    4    4
```

# Layout of several plots on one graph

```
>layout(layout.matrix)
```

